



**Faculty of Natural and
Mathematical Sciences**
Department of Informatics

King's College London
Strand Campus, London,
United Kingdom

7CCSMPRJ

Individual Project Submission 2025/26

Name: Matthew Targett
Student Number: _____
Degree Programme: MSc Computational Finance
Project Title: LLM Proxy-Assisted Combinatorial Auctions: A Modular
Elicitation-Event Framework and Empirical Evaluation
Supervisor: _____
Word Count: 11,948

RELEASE OF PROJECT

Following the submission of your project, the Department would like to make it publicly available via the library electronic resources. You will retain copyright of the project.

Check the appropriate box below

- ☐ I **agree** to the release of my project
☐ I **do not agree** to the release of my project

Signature: _____

Date: August 3, 2026

Abstract

Combinatorial auctions allow bidders to express values for packages of goods, but their formal bid language and exponential bundle space can make participation cognitively and technically demanding. This paper develops a modular framework in which a person describes preferences in natural language, an LLM-based proxy constructs an initial combinatorial bid, and the auction mechanism emits elicitation events that identify economically salient bundles for exact refinement. The framework separates environment generation, preference interpretation, bid construction, mechanism feedback, and evaluation, and is instantiated in iterative sealed and ascending clock auctions. In 95 matched PC-component auction cases, interest-guided candidate generation reduces the number of initially considered bundle values by 55.5% relative to full powerset support. Mechanism-driven refinement raises mean allocation efficiency from 86.4% to 94.8% in the sealed format and 92.8% in the clock format, with the clock using fewer exact value queries. Payment reconstruction remains less accurate than allocation recovery. The results show how natural-language proxy assistance and mechanism feedback can be combined in a reproducible auction architecture, while also identifying the information that sparse elicitation leaves unresolved.

Keywords: combinatorial auctions; preference elicitation; large language models; proxy bidders; XOR bids; VCG; clock auctions; sealed auctions.

Contents

1	Introduction	1
2	Background and Literature Review	2
2.1	Combinatorial Allocation, Incentives, and Communication	2
2.2	Selective Preference Elicitation	2
2.3	Iterative and Clock Auction Design	3
2.4	Machine-Learning-Powered Elicitation	3
2.5	LLM-Based Proxies	3
2.6	Synthesis and Design Requirements	4
3	Formal Framework	4
3.1	Combinatorial Auctions and XOR Bids	4
3.2	Auction Environment and Person Model	5
3.3	Structured Environment Generation	5
3.4	Coverage-Stratified Environment Sampling	6
3.5	Event-Driven Proxy State	6
4	Proxy Inference Pipeline	7
4.1	Natural-Language Questioning	7
4.2	Interest Map	7
4.3	Candidate Bundle Generation	8
4.4	Provisional Valuation and Chunking	8
4.5	Provisional-Value Calibration	9
4.6	Caching, Freezing, and Replay	10
5	Mechanism Environments	10
5.1	Iterative Sealed XOR Auction	10
5.2	Sealed Elicitation Events	10
5.3	Ascending Clock with Supplementary XOR Bids	11
5.4	Clock Elicitation Events	12
5.5	Payments and the Information Required by VCG	12
6	Implementation	13
6.1	Architecture and Information Boundaries	13
6.2	Reliability, Replay, and Auditability	13
6.3	Engineering Challenges and Verification	14
7	Experimental Design	14
7.1	Research Questions	14
7.2	Environments	15
7.3	Experimental Comparisons	15
7.4	Theory-Led Expectations	15
7.5	Metrics	16
7.6	Candidate and Interest-Map Diagnostics	16
7.7	Frozen Specification and Reporting Protocol	17
8	Results	17
8.1	Generated Environment Structure	18
8.2	Provisional-Value Calibration	18
8.3	Interest-Map Support Reduction	18
8.4	Event-Policy Ablations	19

8.5	Full-Suite Overview	19
8.6	Sealed Outcomes and Event Path	20
8.7	Clock Outcomes and Event Path	22
8.8	Cross-Mechanism Comparison and the Payment-Counterfactual Gap	22
9	Discussion	24
9.1	Support Selection and Provisional Scoring	24
9.2	Mechanisms as Elicitation Environments	25
9.3	Non-Monotonic Refinement	25
9.4	Preparation Cost and Person-Side Query Cost	25
9.5	Allocation and Payment Accuracy	25
9.6	Expected and Observed Behaviour	26
10	Legal, Social, Ethical, and Professional Issues	26
10.1	Accessibility, Autonomy, and Distributional Effects	26
10.2	Privacy, Confidentiality, and Security	26
10.3	Economic Trustworthiness and Professional Responsibility	27
11	Limitations	28
12	Extensions	28
13	Conclusion	29
A	Reproducibility and Additional Specification	29
A.1	Frozen Model Roles and Information Boundaries	29
A.2	Frozen Preparation and Mechanism Parameters	29
A.3	Frozen Artifacts and Offline Reconstruction	30
A.4	Software Test Suite	31

List of Figures

1	Implementation architecture and one-way information flow. Hidden values are available to deterministic value-query lookup and evaluation, but not to the person-facing or bidder-proxy model calls.	13
2	Structure of the generated experimental environment. Panel (a) reports positive bidders per master-population good; panel (b) reports bidder interest density by stratum; panel (c) reports full-information winner counts across accepted cells, with colour indicating the largest winner’s welfare share and marker size indicating bidder count.	18
3	Auction-level candidate-support size with eight bidders as the number of goods varies. Points are five-seed means; the vertical axis is logarithmic. The full baseline is $8(2^{ G } - 1)$	19
4	Five-seed event-policy ablations at the 8×8 anchor.	20
5	True allocation efficiency across the bidder, goods, and joint scalability paths. Lines show cell means over five seeds and points show seed-level observations. The 8×8 anchor appears in each panel.	21
6	Deterministic exact value queries across the final scalability suite. Offline disclosure, interest-map, and provisional-valuation calls are excluded.	21
7	Reported VCG-witness incidence by event type over the final 95-case suites. Incidence is descriptive rather than a causal ablation.	22
8	Illustrative final-suite trajectories: the largest sealed efficiency gain (left) and largest clock terminal welfare reduction (right). These are diagnostic examples, not representative averages.	23
9	Paired sealed and clock efficiency under the same frozen initial proxy state. The dashed line denotes equal efficiency.	23
10	Absolute reported-space VCG payment error normalised by full-information optimum welfare. Lower is better.	24

List of Tables


1	Relationship between prior elicitation approaches and this project.	4
2	Feedback signals by mechanism.	13
3	Theory-led expectations and their empirical tests.	16
4	Information and cost layers in the experimental architecture.	17
5	Final allocation and communication results over 95 matched cases. Efficiency intervals use five seed-level means. Better/same/worse compares each elicited mechanism with its common initial provisional allocation.	20
6	Different information requirements for allocation and VCG pricing.	24
7	Principal frozen parameters. Zero query caps denote disabled caps.	30

1 Introduction

Combinatorial auctions allocate heterogeneous goods when values depend on packages rather than isolated items. They can express complements, substitutes, and alternative configurations, but that expressiveness creates two barriers to participation. First, m goods induce 2^m possible bundles. Second, bidders must translate informal preferences into a formal bid language and reason about package values. Even where computation is available, this reporting task can impose substantial cognitive and technical burden.

LLM-based proxies offer a possible interface between a person and that formal market. A bidder can communicate in ordinary language, while a proxy interprets the disclosure, constructs a combinatorial bid, and responds to structured feedback [Huang et al., 2025]. Such assistance could make expressive auctions accessible to participants who would otherwise struggle to formulate XOR package bids. This paper does not conduct a human-subject study and therefore does not claim measured reductions in cognitive burden. Instead, it develops and tests the computational architecture required for that interface.

The central proposal is an *elicitation-event framework*. A proxy first converts natural-language evidence into a structured interest map and an initial set of provisionally valued bundles. During the auction, the mechanism emits events that expose allocation-relevant or payment-relevant uncertainty. The proxy maps each eligible event to a formal refinement action. This makes the mechanism an elicitation environment rather than merely a passive consumer of a one-shot LLM bid. The organising information path is



Two contrasting mechanism families are evaluated:

1. an **iterative sealed XOR auction** with provisional feedback and final VCG payments;
2. an **ascending clock mechanism** with supplementary XOR atoms and final VCG payments.

They share the same initial proxy state and formal bid representation but expose different evidence. The sealed mechanism repeatedly reveals provisional allocations and counterfactual competitors. The clock produces a price and demand trajectory that restricts the terminal verification frontier. Studying both formats tests whether the event abstraction can accommodate distinct mechanism-specific routes to economically salient bundles.

The event interface is designed to be extensible to other mechanisms, but the empirical evidence is narrower: it covers two implemented mechanism families, one generated PC-component master population, and one frozen model configuration. The paper therefore distinguishes architectural generality from empirical generalisation throughout.

The main research question is:

How can an LLM proxy translate natural-language preferences into a tractable combinatorial bid, and how effectively can mechanism-specific elicitation events refine that bid toward efficient allocation?

The primary contribution is the modular elicitation-event interface. Three supporting contributions make that interface testable and operational:

1. **A modular elicitation-event interface.** Sealed and clock mechanisms expose different economic signals through a common event abstraction, allowing targeted, deduplicated refinements and paired evaluation from the same frozen initial state.
2. **Interest-guided candidate support and provisional initialisation.** A person-specific interest map determines what the proxy considers; provisional values then rank that support. This separates omission from scoring error.

3. **Structured test environments.** A constrained pipeline produces heterogeneous private preferences and deterministic full-information benchmarks for auditable evaluation.
4. **Information-aware diagnostics.** Candidate coverage, allocation trajectories, event incidence, and VCG witness logging identify whether information is lost during support selection, provisional scoring, allocation refinement, or payment-counterfactual construction.

Caching, chunked structured output, fail-closed parsing, frozen replay, and detailed provenance are supporting experimental controls rather than independent conceptual contributions.

The report proceeds from related auction and human–AI work to the formal proxy model and inference pipeline. It then defines the sealed and clock event environments, describes the implementation and experimental protocol, and evaluates support reduction, policy selection, scalability, allocation, and payment recovery. The discussion separates supported findings from limitations, and Section 10 considers the safeguards that real bidder-facing deployment would require.

2 Background and Literature Review

This project lies at the intersection of mechanism design, selective preference elicitation, machine-learning-powered auctions, and human–AI interaction. The relevant literature supplies two different kinds of constraint. Auction theory identifies the information needed for allocation and payment; elicitation and interaction research asks how that information can be acquired without requiring a person to enumerate an exponential bid.

2.1 Combinatorial Allocation, Incentives, and Communication

Combinatorial auctions are appropriate when values contain complements or substitutes, but unrestricted valuations contain one number for every bundle. Nisan and Segal show that efficient allocation can require exponential communication for general valuation classes [Nisan and Segal, 2006]. This is not merely a solver problem: no interface can guarantee exact general preference recovery from a small transcript. Restricted bid languages, structured valuation classes, approximation, and selective elicitation are therefore substantive assumptions rather than implementation conveniences.

Vickrey and Groves mechanisms establish the benchmark for efficient truthful allocation under quasilinear private values [Vickrey, 1961, Groves, 1973]. In a combinatorial setting, however, a winner’s VCG payment depends on welfare in a bidder-removal economy. Accurate allocation and accurate payment thus require overlapping but non-identical information. This distinction motivates the paper’s separate allocation and VCG-witness diagnostics. It also limits deployment claims: sparse reported bids can reproduce the efficient allocation while leaving counterfactual payments inaccurate.

2.2 Selective Preference Elicitation

Lahaie and Parkes connect preference elicitation to computational learning theory. Value queries can simulate membership queries, while some demand-style queries resemble equivalence-query counterexamples [Lahaie and Parkes, 2004]. The analysis supplies guarantees for particular valuation classes and bid languages, but assumes that bidders can answer the formal queries they receive. In practice, choosing a maximum value for an arbitrary bundle or solving a demand problem at posted prices may itself be a substantial cognitive task.

Work on structured classes such as strong substitutes illustrates the same tension. Goldberg, Lock, and Marmolejo-Cossío study learning strong-substitutes demand via queries and show efficient algorithms for restricted settings but hard lower bounds in more general cases [Goldberg

et al., 2022]. The theoretical lesson is that structure matters. The practical question is how to extract useful structure from human preference descriptions without requiring full formal revelation.

2.3 Iterative and Clock Auction Design

Iterative combinatorial auctions reduce upfront reporting by eliciting information over rounds. Parkes surveys mechanisms that alternate among provisional allocations, prices, and new bids [Parkes, 2006]. iBundle uses ascending package prices to direct bidders towards welfare-relevant packages [Parkes and Ungar, 2000]; ascending proxy auctions similarly use bidder-side proxy agents to express package demand [Ausubel and Milgrom, 2006]. These mechanisms motivate the sealed arm’s use of repeated WDP feedback, but they do not by themselves resolve how a non-expert person communicates values to a proxy.

Clock auctions expose a different information structure: item prices generate a demand path before supplementary package information is considered. The clock-proxy design combines transparent price discovery with a later package phase intended to recover efficiency [Ausubel et al., 2006]. Core-selecting payments are often used in package-auction practice to address properties of VCG outcomes and incentives [Day and Milgrom, 2008]. The present clock is therefore CCA-inspired rather than a production CCA: its purpose is to test whether the revealed-demand path can serve as a relevance filter for later exact verification.

That simplification creates a deliberate trade-off. Using VCG in both arms makes reported-payment error comparable, while omitting activity constraints, optimised increments, and core-selecting payments prevents the clock results from being interpreted as evidence about a deployable spectrum-auction design.

2.4 Machine-Learning-Powered Elicitation

Machine-learning-powered combinatorial auctions use partial bidder responses to train surrogate valuation models and select future queries. MLCA uses learned models to guide value-query selection and reports high efficiency with comparatively few elicited values [Brero et al., 2019]. iMLCA adds interval bids and activity rules, weakening the requirement that every answer be an exact value [Lubin et al., 2020]. These approaches learn a numerical model from formal responses; the current project instead begins with qualitative language and separates support inference from numerical scoring.

More recent ML-powered clock mechanisms learn from demand information and select queries with high clearing potential [Soumalias et al., 2023]; hybrid work combines prices, bids, and values [Soumalias et al., 2024]. This literature establishes welfare-versus-query-cost analysis as a natural standard and shows that the mechanism’s observable state should influence query selection. It also provides the principal alternative to the event framework: a learned acquisition rule could choose queries directly. The deterministic event policies used here sacrifice such adaptivity for auditability and exact paired replay.

2.5 LLM-Based Proxies

Huang et al. propose LLM-based proxies for accelerated preference elicitation in combinatorial auctions [Huang et al., 2025]. Their proxies maintain state and combine natural-language questions with formal value and demand queries. Soumalias et al. likewise study LLM-powered elicitation in combinatorial assignment [Soumalias et al., 2025]. These studies show the promise of language as an interface, but the numerical reliability of LLM answers and the mechanism-specific nature of the interaction remain design concerns.

Human-AI interaction research adds requirements that auction metrics alone do not capture. Guidance synthesised by Amershi et al. [2019] emphasises setting expectations, exposing un-

certainty, supporting correction, and giving people control when an AI system is wrong. In an auction, an inference error can affect payments and other bidders as well as interface satisfaction. The current frozen experiment does not test those human outcomes, but its explicit interest map, query log, and deterministic correction path are steps towards an inspectable rather than opaque proxy.

2.6 Synthesis and Design Requirements

Table 1 locates the project relative to the main approaches. The gap addressed is not a claim to solve unrestricted preference recovery. It is the absence of a reproducible interface that separates language-based support discovery, provisional numerical scoring, and mechanism-specific exact refinement while allowing different auction formats to share the same initial proxy state.

Table 1: Relationship between prior elicitation approaches and this project.

Approach		Main evidence	Strength	Boundary motivating this work
Classical elicitation	selective	Formal value or demand queries	Theoretical query guarantees for structured values	Assumes a bidder can answer the formal query
ML-powered auctions	auc-	Values, intervals, or demands	Learns where further information is useful	Numerical model and acquisition policy are often mechanism-specific
Clock/proxy auctions	auc-	Ascending prices and package demand	Price discovery restricts relevant outcomes	Does not supply a natural-language person-proxy interface
LLM proxy elicitation	elicita-	Language plus formal queries	Translates qualitative goals into bids	Model noise and protocol portability remain concerns
This project		Validated disclosure, inferred support, provisional values, and event-selected exact VQs	Paired, auditable support-score-refine pipeline	Frozen non-agentic policy and one generated domain

The synthesis yields five requirements used in the design below: preserve a hidden full-information benchmark; distinguish candidate omission from value misranking; prevent model calls from accessing hidden values; expose a common typed event boundary while retaining mechanism-specific triggers; and freeze initial inference so that sealed and clock outcomes are paired. These choices trade open-ended proxy autonomy for experimental identification.

3 Formal Framework

3.1 Combinatorial Auctions and XOR Bids

Let $G = \{g_1, \dots, g_m\}$ be a finite set of indivisible goods, and let $\mathcal{B} = 2^G$ be the bundle space. There are bidders $N = \{1, \dots, n\}$. Bidder i has a hidden valuation $v_i : \mathcal{B} \rightarrow \mathbb{R}_{\geq 0}$. An allocation is $a = (a_1, \dots, a_n)$, where $a_i \subseteq G$, subject to feasibility:

$$a_i \cap a_j = \emptyset \quad \forall i \neq j.$$

The full-information efficient allocation is

$$a^* \in \arg \max_a \sum_{i \in N} v_i(a_i).$$

The implementation represents formal bids using XOR atoms. An XOR bid for bidder i is a finite set

$$\theta_i = \{(b, \hat{v}_i(b)) : b \in B_i\},$$

where $B_i \subseteq 2^G$. The induced reported valuation is

$$v^{\theta_i}(b) = \max_{\substack{b' \in B_i \\ b' \subseteq b}} \hat{v}_i(b'),$$

with value zero if no reported atom is contained in b . This free-disposal interpretation is used by the WDP solver.

Given reported bids $\theta = (\theta_1, \dots, \theta_n)$, the WDP is

$$\max_a \sum_i v^{\theta_i}(a_i)$$

subject to feasibility. The implementation solves this problem by an integer linear programming formulation over XOR atoms.

3.2 Auction Environment and Person Model

An auction environment is

$$\mathcal{E} = (G, N, D_A, D_G, (\Lambda_i)_{i \in N}, (v_i)_{i \in N}),$$

where D_A is the auction description, D_G maps goods to public text descriptions, Λ_i is bidder i 's private structured preference profile, and v_i is the deterministic hidden benchmark valuation induced by that profile. The bidder proxy observes the public catalogue and the person's disclosed answer, but never observes Λ_i or v_i . The person simulator and evaluator may access the private profile for different purposes.

The person simulator is distinct from the bidder proxy. Before an experiment, the profile is rendered deterministically into a brief qualitative disclosure contract D_i . A person model conditioned on D_i answers one canonical opening question, disclosing natural-language goals, positively valued options, relevant alternatives or complements, acquisition mode, explicit exclusions, and an approximate total budget. The answer is independently verified against Λ_i before it is accepted.

During the principal mechanism runs, a value query for bundle b is answered deterministically as $v_i(b)$. This removes repeated LLM response noise and isolates the bundle-selection policy induced by the auction mechanism. LLM-based value and demand answers remain available as robustness treatments, but they are not used in the main scalability specification.

The implementation records three model roles separately: environment, person, and bidder proxy. The environment and person roles may use the same model, while the bidder proxy can use a different provider and model. In the current frozen population the environment and person roles use `gpt-5.6-sol`, while interest-map and provisional-valuation inference use `gemini-3.6-flash`. Role separation is an information boundary, not necessarily a requirement that all roles use different model families.

3.3 Structured Environment Generation

The environment pipeline begins from a human-specified population design, not an unconstrained request for an auction. The design fixes the 16 good identifiers and descriptions, their functional categories, 16 bidder archetypes across four broad strata, and population- and sample-level constraints. It deliberately leaves bidder-specific numeric and structural preferences to the environment model.

Each bidder is generated from a latent profile

$$\Lambda_i = (w_i, \mathcal{S}_i, \mathcal{C}_i, B_i, \rho_i, D_i),$$

where w_i gives base item values, $\mathcal{S}_i$ person-specific substitute groups, $\mathcal{C}_i$ complement groups, B_i a budget cap and scale, ρ_i saturation behaviour, and D_i a role and identity description. Every good is also classified as core, secondary, low interest, or zero value for that bidder.

For a substitute group $S \in \mathcal{S}_i$, the profile gives full value to the highest-valued selected item and discounted backup or resale value to additional selected items:

$$\text{sub}_i(S, b) = w_i(g^*) + \lambda_{iS} \sum_{g \in (b \cap S) \setminus \{g^*\}} w_i(g), \quad g^* \in \arg \max_{g \in b \cap S} w_i(g).$$

A complement group $C \in \mathcal{C}_i$ contributes a bonus α_{iC} when $C \subseteq b$. For ordinary single-system buyers, $\lambda_{iS} = 0$ implements a **choose_one** relationship. For resellers, integrators, or multi-machine buyers, a positive backup factor implements **can_use_multiple**. Thus functional similarity does not imply that joint ownership is valueless. The generated value is a sum of ordinary item contributions, substitute-group contributions, complement bonuses, and saturation or cap adjustments. Values are repaired to satisfy free disposal.

The environment model generates the structured profile rather than directly enumerating a value for every bundle. Deterministic code then stores complete hidden values for every non-empty bundle. For m goods this gives $2^m - 1$ benchmark values per bidder. These tables support full-information allocation, exact value queries, welfare evaluation, and diagnostic decomposition.

Generation proceeds in batches with bounded automatic repair. A batch is accepted only if its profiles satisfy schema, classification, monetary, and person-level interest-density constraints. The assembled population is then validated for per-good positive and core bidder coverage, cross-stratum competition, multiway substitute structure, complement prevalence, and acquisition-mode consistency.

3.4 Coverage-Stratified Environment Sampling

Experiments use a fixed validated 16-by-16 master population but vary the selected goods and bidders. For each seed, a deterministic coverage-stratified search creates one nested ordering of all goods and one nested ordering of all bidders. Prefixes of these latent orderings generate the required sizes while the public scenario preserves catalogue declaration order.

Every scalability composition is validated for minimum interest coverage, positive value for each selected bidder, complement and substitute prevalence, category and bidder-stratum diversity, at least two full-information winners, and a bound on the largest winner’s welfare share. The five selected composition seeds produce distinct good sets, bidder sets, and combined environments for every experimental cell. They are alternative nested orderings of one master population, not independently generated populations.

3.5 Event-Driven Proxy State

Bidder i ’s proxy has state

$$S_i = (K_i, I_i, \mathcal{C}_i, O_i, \tilde{v}_i, \theta_i),$$

where K_i is the knowledge base, I_i the interest map, $\mathcal{C}_i$ the candidate bundle support, O_i explicit observations, $\tilde{v}_i$ provisional or refined value estimates, and θ_i the mechanism-facing XOR bid.

An elicitation event is represented as

$$e = (\tau, i, b, x, \rho),$$

where τ is the event type, i the target bidder, b an optional target bundle, x the mechanism state that triggered the event, and ρ its eligibility evidence or economic rationale. This representation separates three operations:

$$E_t = \text{Emit}_M(x_t), \quad q_t = \text{Select}_\pi(E_t, S_t), \quad S_{t+1} = \text{Update}(S_t, q_t, y_t),$$

where the mechanism emits eligible events, the event policy selects an action, and the proxy updates its state after response y_t . A policy may prioritise or ignore emitted events without changing the auction implementation, while a new mechanism can emit the same event type from different state evidence.

The proxy responds to elicitation events:

$$\pi_i : \mathcal{E}_{event} \times S_i \rightarrow (\mathcal{A}_{proxy}, S'_i).$$

As a general interface, events may trigger language interaction, structural inference, value or demand refinement, bid submission, or no action. This expressiveness should not be confused with the experimental instantiation in this paper. Here the opening exchange, interest map, candidate support, and provisional values are frozen before the auction. At auction time, a deterministic event policy selects an unqueried bidder–bundle pair, exact value lookup deterministically returns the hidden benchmark value, and that value replaces the provisional report. The runner does not deliberate over a transcript, autonomously choose a new query type, or make any auction-stage LLM call. The proxy state therefore defines a reusable general architecture, while the experiments evaluate a deliberately non-agentic event-policy instantiation.

The generic update is

$$e_t \rightarrow \pi_i(e_t, S_{i,t}) \rightarrow q_{i,t}^V(b) \rightarrow v_i(b) \rightarrow S_{i,t+1} \rightarrow m_{i,t+1}.$$

For sealed auctions, $m_{i,t+1}$ is an XOR bid. For clock auctions, it is a demand response and supplementary XOR information.

4 Proxy Inference Pipeline

4.1 Natural-Language Questioning

The current main configuration uses one canonical opening question tailored to the PC-component catalogue. It asks what the person hopes to obtain, how the alternatives fit their needs, whether any goods are useful only together or remain useful in multiples, and approximately how much they would spend overall. The same question is used across bidders in a given catalogue so that differences in their answers arise from preferences rather than question wording.

The person answer is intentionally brief and qualitative. It is not intended to recover a complete value table. It discloses enough information to support candidate selection and approximate scale without exposing base values, complement bonuses, or the hidden bundle table. The answer and its preparation-time verification record are frozen before either auction mechanism is run. Proxy-generated opening questions remain an optional robustness treatment.

4.2 Interest Map

The proxy derives a lightweight interest map from the auction description, item descriptions, natural-language question, and person answer:

$$I_i = (G_i^+, G_i^-, C_i, S_i, \beta_i, r_i).$$

Here G_i^+ is the inferred interested item set, G_i^- excluded items, C_i complement groups, S_i substitute groups, β_i a budget hint, and r_i a rationale.

Each inferred substitute group additionally records an acquisition mode

$$\mu(S) \in \{\text{choose_one}, \text{can_use_multiple}, \text{unclear}\}$$

and textual evidence. A restrictive `choose_one` mode is accepted only when the answer states that at most one member is useful or that additional members add no meaningful value. Ranking one good above another, calling an item a fallback, or observing functional similarity is insufficient. This distinction prevents a GPU category, for example, from being incorrectly treated as exclusive for a reseller who can use multiple cards.

The main implementation uses a closed-world map: every catalogue item must be classified, and an item that is neither positively disclosed nor included in a disclosed structural group is placed in G_i^- . Candidate generation considers subsets of interested items, never includes excluded items, and filters joint bundles only when they contain multiple members of an evidence-grounded `choose_one` group. `Can_use_multiple` and `unclear` groups are retained conservatively.

The implementation enforces only generic consistency. It does not impose hidden PC-build-specific rules or use the private profile to repair the proxy’s map. Hidden truth is used solely for diagnostics: item precision/recall, exact group precision/recall, acquisition-mode accuracy, choose-one precision/recall, dangerous false exclusivity, and oracle-versus-inferred candidate-set precision and recall.

The current implementation also avoids silent full-powerset fallback in live experiments. If interest-map parsing fails after retries, the run can fail closed under an interest-map failure policy. A degraded all-items fallback remains available for debugging, but it is not the preferred experimental setting because it can generate excessively broad and incoherent support.

4.3 Candidate Bundle Generation

Given I_i , candidate generation produces a priority-ordered set

$$\mathcal{C}_i = \Psi(I_i, G, M_{\text{cand}}).$$

The implemented ordering is:

1. complete declared complement groups;
2. singletons of interested items;
3. all remaining valid subsets in ascending size and lexicographic order.

A bundle is valid if it does not include multiple goods from the same declared `choose_one` substitute group. Bundles containing multiple members of `can_use_multiple` or `unclear` groups remain eligible. A candidate cap M_{cand} can be applied after prioritisation, although the principal scalability preparation uses $M_{\text{cand}} = \infty$ and therefore leaves the inferred support uncapped.

This design intentionally keeps the interest map simple. Its main function is to reduce support size while preserving important complement groups and singletons. Failure modes are logged through candidate-count diagnostics, candidate reduction ratios, coverage measures, and interest-map quality flags such as broad interest sets with no substitute groups.

4.4 Provisional Valuation and Chunking

After candidate generation, the proxy assigns provisional values:

$$\tilde{v}_i(b), \quad b \in \mathcal{C}_i.$$

The proxy model receives the public catalogue, opening exchange, interest map, and enumerated candidate support. It returns one estimated value for each candidate bundle. These values both

rank the support and initialise the proxy’s mechanism-facing XOR bid. They are proxy-side inferences, not person value queries, and they do not have access to the private profile or hidden valuation table.

For large supports, the implementation uses deterministic PV chunking. If

$$|\mathcal{C}_i| > K_{PV},$$

candidate bundles are split into ordered chunks of at most K_{PV} bundles. Each chunk is valued by a separate provisional-valuation call and the results are merged. Chunking preserves the semantics of the initial bulk PV stage while reducing the risk of truncation or malformed large JSON outputs. PV chunks are counted as shared initial proxy inference calls, not value-query refinements.

The current preparation configuration uses chunks of at most 64 bundles and fail-closed PV behaviour. If a chunk fails after parse retries, the run aborts with a diagnostic error identifying the bidder, candidate count, chunk size, token budget, and original parsing exception. A degraded zero-initialisation policy exists for debugging, but failed PV never triggers mass direct value queries over the candidate support. Otherwise a structured output failure would silently change the elicitation treatment and contaminate person-query counts.

4.5 Provisional-Value Calibration

Provisional valuations are estimated from the person’s short qualitative disclosure alone, and the estimator’s error need not be centred. The implementation therefore supports an explicit calibration of the raw provisional value, parameterised by a family, a multiplicative scale, and an optional bundle-size adjustment:

$$\hat{v}^{cal}(b) = \min \left\{ \bar{B}, s \cdot \hat{v}^{raw}(b) \cdot \gamma^{\max(0, |b| - k_0)} \right\}.$$

Here $s > 0$ is the scale, $\gamma > 0$ the size factor, k_0 the size threshold, and $\bar{B}$ the bidder’s *disclosed* budget, taken from the elicited interest map rather than from any hidden value. The scale is not restricted to $(0, 1]$: the estimator’s qualitative conservatism instructions mean the observed bias is typically downward, so the corrective scale is generally greater than one. The family selects which terms apply: the identity (raw values), scale only, or scale with the size adjustment.

Calibration applies only to initial LLM-inferred provisional values. Exact deterministic value-query answers and mechanism-triggered refinements are ground truth and are reported unchanged.

Frozen elicitation packs store raw, uncalibrated provisional values, so alternative calibration rules are replayed over an identical elicitation without new LLM calls; a calibrated and an uncalibrated arm therefore differ only in the calibration and not in re-elicitation noise. Every run records its fully resolved effective calibration, together with a stable configuration hash, in its console header, machine-readable run configuration, and result tables.

The final specification uses the accepted uniform scale $s = 1.4758$, a disclosed-budget cap, and no bundle-size adjustment ($\gamma = 1$). It was selected before the main mechanism runs using nine synthetic consumer-bundle environments spanning camera/video, kitchen, and travel domains, with three independently generated instances per domain. Three folds leave out one instance index across all domains, fitting on six environments and evaluating on the remaining three. The objective is mean absolute reported-VCG payment error normalised by optimum welfare, subject to a mean-efficiency tolerance. This procedure shares no goods or bidder profiles with the PC-component population. On the nine held-out evaluations, calibration reduces the objective from 0.289 to 0.233 and improves seven of nine cases. The accepted calibration and its hash are frozen in the final experiment specification.

Algorithm 1 Natural-language disclosure to initial XOR bid

Require: Public catalogue G , canonical question, person disclosure contract D_i

- 1: Obtain and independently verify opening answer A_i
 - 2: Infer interest map I_i from (G, A_i)
 - 3: Generate $\mathcal{C}_i = \Psi(I_i, G, M_{\text{cand}})$
 - 4: Estimate $\tilde{v}_i(b)$ for every $b \in \mathcal{C}_i$, using deterministic chunks if required
 - 5: Apply the frozen calibration rule and disclosed-budget cap
 - 6: **return** $\theta_i^0 = \{(b, \tilde{v}_i(b)) : b \in \mathcal{C}_i\}$ and frozen proxy state S_i^0
-

4.6 Caching, Freezing, and Replay

LLM calls are cached by a hash of the rendered prompt and all response-relevant metadata. Cache modes include off, read-write, read-only, and refresh. Each successful response is committed independently, allowing interrupted preparation to resume without repeating completed calls.

Caching supports preparation, but the principal experimental control is the frozen elicitation pack. A pack contains the accepted opening exchange, interest map, candidate bundles, raw provisional values, model provenance, generation settings, and call records for one catalogue and bidder set. Fresh mutable proxies replay this immutable initial state with zero LLM calls. Consequently sealed and clock arms differ only in their mechanism paths and deterministic refinement queries.

Initial elicitation depends on the selected goods but not on the number of rival bidders. For each seed, preparation therefore generates one master pack for every goods count and then projects ordered bidder subsets into the 19 scalability cells. This avoids regenerating the same bidder-catalogue exchange for the fixed-goods bidder sweep.

5 Mechanism Environments

5.1 Iterative Sealed XOR Auction

At round r , each proxy submits an XOR bid θ_i^r , and the mechanism solves:

$$a^r \in \arg \max_a \sum_i v_i^{\theta_i^r}(a_i).$$

The WDP and reported allocation provide the mechanism state from which events are emitted; the auction algorithm itself does not prescribe which bundle is queried. After each set of event responses, corrected XOR atoms replace their provisional counterparts and the WDP is solved again. The main termination condition is the absence of any eligible new refinement after at least one cycle. Forty rounds is a non-binding runaway safeguard in the final suite. Every bidder–bundle query is deduplicated globally.

5.2 Sealed Elicitation Events

Four event families are considered because they expose different parts of the reported allocation frontier.

Incumbent verification. The provisionally allocated bundle is the atom whose error can most directly invalidate the reported winner. Whenever a bidder receives a non-empty, previously unqueried bundle in the current WDP allocation, the event targets that bidder–bundle pair for exact value lookup.

Winner-removal allocation counterfactual. VCG payments and competitive displacement depend on allocations that become feasible when a current winner is absent. For each winning atom, the mechanism temporarily removes that atom, re-solves the WDP, and emits events for previously unqueried bundles that enter the counterfactual allocation.

Scarcity-avoiding fallback. A high provisional incumbent may occupy contested goods even when its bidder has a credible alternative. If an incumbent contains currently contested goods, the event selects the bidder’s highest-reported unqueried alternative that reduces contested-good exposure and retains at least 60% of the incumbent’s reported value. It is triggered at most once per bidder.

Large-correction follow-up. A large initial error is evidence that a nearby provisional ranking may also be unreliable. If exact lookup changes a value by at least 25% under the symmetric relative correction measure, the event targets one unqueried structurally neighbouring candidate bundle. It is triggered at most once per bidder.

Independent loser exploration is disabled in the principal experiments: losers are queried only when they enter a winner-removal allocation. Section 8.4 selects among combinations of the four events; the definitions above do not presuppose that every event is empirically useful.

Algorithm 2 Iterative sealed elicitation

Require: Frozen initial bids θ^0 , sealed event policy π_S

```

1: for round  $r = 0, 1, \dots$  do
2:   Solve the reported WDP to obtain  $a^r$ 
3:    $E_r \leftarrow \text{Emit}_S(a^r, \theta^r)$ 
4:   Remove previously queried bidder–bundle pairs from  $E_r$ 
5:    $Q_r \leftarrow \text{Select}_{\pi_S}(E_r, S^r)$ 
6:   if  $Q_r = \emptyset$  after at least one elicitation cycle then
7:     break
8:   end if
9:   for all  $(i, b) \in Q_r$  do
10:    Replace  $\tilde{v}_i(b)$  by exact  $v_i(b)$ 
11:   end for
12: end for
13: return final reported allocation and reported-space VCG payments

```

5.3 Ascending Clock with Supplementary XOR Bids

The clock arm posts item prices p_t . The reported surplus of bundle b for bidder i is

$$u_{i,t}^\theta(b) = v^{\theta_{i,t}}(b) - \sum_{g \in b} p_{t,g}.$$

Each proxy reports up to k top-surplus bundles; the principal specification uses $k = 3$. The mechanism increments every good with positive excess demand by \$50. All top- k demands observed along the price path are accumulated as supplementary XOR atoms. No exact value query is made while prices ascend: price discovery is therefore allowed to complete on the frozen calibrated reports. The clock stops when there is no excess demand; 500 rounds is a runaway safeguard rather than a target.

Thus price discovery is query-free and naturally terminates before terminal exact refinement begins. Let R_i denote bidder i ’s bundles revealed among its top- k demands. These revealed sets delimit the clock-specific event space: unrevealed losing bundles do not become eligible merely because they have high provisional values.

5.4 Clock Elicitation Events

Single-pass revealed VCG witness. Before any correction, the mechanism removes each provisional winner and solves the associated supplementary counterfactual allocation. This freezes a pre-query witness frontier. An event is emitted only when a counterfactual bundle is both unqueried and belongs to its bidder’s revealed-demand set R_i . The event targets payment-relevant competitors while preserving the support discipline supplied by clock price discovery.

Winner closure. Corrections can change the supplementary allocation. After each re-solve, an event is emitted for every newly allocated, previously unqueried bundle. The mechanism alternates exact lookup and re-optimisation until all current winners have been verified. This targets allocation correctness rather than broad frontier exploration.

Staged revealed-witness closure. Winner corrections can expose new bidder-removal competitors. The mechanism therefore recomputes the revealed witness frontier after winner closure, queries eligible revealed witnesses, and returns to winner closure. The two stages alternate until neither produces a new query. Section 8.4 evaluates this sequence against narrower and broader terminal policies.

Algorithm 3 Clock price discovery and terminal revealed-winner verification

Require: Frozen bids θ^0 , top- k rule, price increment Δp

- 1: **while** some good has positive excess demand **do**
 - 2: Record each bidder’s top- k positive-surplus demands in R_i
 - 3: Raise each over-demanded good’s price by Δp
 - 4: **end while**
 - 5: Construct supplementary XOR bids from all revealed demands
 - 6: Freeze and query eligible pre-correction revealed VCG witnesses
 - 7: **repeat**
 - 8: Query newly allocated unverified bundles until winner closure
 - 9: Recompute and query eligible revealed bidder-removal witnesses
 - 10: **until** neither stage emits an unqueried bidder–bundle pair
 - 11: **return** final supplementary allocation and reported-space VCG payments
-

5.5 Payments and the Information Required by VCG

VCG is retained as a common reported-space payment rule so that sealed and clock allocation paths remain comparable. Payment accuracy is analytically distinct from allocation efficiency. Even if every finally allocated bundle has been queried exactly and the efficient allocation has been found, bidder i ’s payment depends on the welfare attainable by the other bidders when i is removed. Sparse provisional reports can therefore yield accurate allocations but inaccurate winner-removal counterfactual welfare and revenue.

Allocation efficiency is consequently the primary mechanism outcome. Reported VCG revenue, full-information VCG revenue, and per-winner payment error are reported as secondary diagnostics. A post-allocation payment-counterfactual audit is a natural extension, but replacing VCG with pay-as-bid would conflate preference elicitation with strategic bid shading and is not part of the principal design.

Table 2: Feedback signals by mechanism.

Mechanism	Main feedback	Proxy refinement opportunity
Sealed XOR-VCG	Provisional WDP, allocated bundles, winner-removal frontier, contested goods	Refine current winners, counterfactuals, scarcity fallbacks, and one-step correction neighbours
Clock with supplementary VCG	Ascending prices, top-three revealed demands, terminal winner-removal frontier	After price discovery, close newly winning bundles and revealed VCG witnesses

6 Implementation

6.1 Architecture and Information Boundaries

The implementation separates preparation-time model orchestration from deterministic auction execution. Figure 1 shows the principal data path. Environment generation produces private structured profiles and deterministic full valuation tables. A person model sees only a profile-specific disclosure contract; the bidder-proxy model sees the public catalogue and accepted opening exchange. The resulting interest map, candidate support, and raw provisional values are frozen. Mechanism runners then replay that state without making any model call.




Figure 1: Implementation architecture and one-way information flow. Hidden values are available to deterministic value-query lookup and evaluation, but not to the person-facing or bidder-proxy model calls.

Pure mechanism modules receive typed bids and proxy states, solve the XOR WDP, and emit typed elicitation events. They do not construct prompts or know which model generated the initial state. The event-policy layer converts eligible mechanism evidence into a bidder-bundle refinement request. The proxy adapter deduplicates that request, performs an exact hidden-value lookup in the principal experiment, replaces the corresponding provisional atom, and returns the revised formal bid. This boundary is what permits the same proxy representation to support two mechanism-specific policies.

6.2 Reliability, Replay, and Auditability

Preparation uses strict schemas and fail-closed validation. A malformed or semantically inconsistent structured response is retried with diagnostic context and then aborts; the principal pipeline does not silently substitute an all-items interest map or a mass of artificial zero values. Provisional values are chunked because a single response over a large candidate support was both slow and prone to truncation. Chunk outputs are validated against the requested bundle identifiers before they are merged.

Long-running preparation is interruption-safe. Each accepted response is committed to a SQLite cache keyed by the rendered prompt, provider, model, and response-relevant settings. Completed goods-specific master packs can be resumed and projected to smaller bidder sets without regenerating responses. Every frozen pack includes an environment fingerprint and is rejected if replayed against a different catalogue or bidder composition. Raw and parsed responses, token counts, cache status, model roles, and configuration hashes remain available for audit, while the public serialisation omits the complete hidden valuation table.

6.3 Engineering Challenges and Verification

Several design features arose from observed failures rather than from the initial architecture alone. Generated profiles sometimes omitted goods or placed zero-valued goods in substitute groups; schema checks, deterministic repairs, and population-level economic preflight now prevent such profiles from entering the experiment. Broad bidder interest caused candidate growth and oversized model responses; person-specific exclusion, evidence-grounded `choose_one` filtering, and chunking address different parts of that problem. Repeated mechanism signals initially risked asking the same exact question more than once; deduplication is therefore enforced on the bidder–bundle pair independently of event type.

Clock experiments also showed why termination must be distinguished from a runtime safeguard. A 50-round limit truncated legitimate low-increment price paths, so the final design stops on zero excess demand and uses 500 rounds only as a fault safeguard. Finally, high allocation efficiency initially obscured large payment errors. Explicit winner-removal witness logging was added so that a refinement can be classified by both allocation incidence and reported-VCG-counterfactual incidence.

The repository’s offline suite contains 880 tests spanning these boundaries: WDP and payment invariants, mechanism trajectories, event deduplication, interest-map filtering, parsing and failure policies, caching, frozen-pack integrity, calibration, scalability grids, and final input hashes. Mock clients ensure that tests require no credentials or live calls. The suite and its limitations are described in Appendix A.4. Model identifiers, token limits, frozen policy parameters, artifact paths, and reconstruction commands are collected in Appendix A.

7 Experimental Design

7.1 Research Questions

The empirical design addresses:

1. Does the structured generation pipeline produce diverse, competitive, and economically non-trivial environments across the scalability grid?
2. How much does interest-map inference reduce candidate support relative to full powerset reporting?
3. How much welfare do calibrated provisional bids recover before exact refinement?
4. Which sealed and clock event policies improve allocation efficiency, and at what exact-query cost?
5. How do support, allocation efficiency, and exact-query burden scale with goods and bidders?
6. Why can reported-space VCG payments remain inaccurate when the allocation is efficient, and which event types reach payment witnesses?

7.2 Environments

The master population contains 16 PC-component goods and 16 bidders in four strata: gaming/performance, budget/office, professional/AI, and reseller/procurement. The bidder set includes enthusiast and value gamers, a streamer, an overclocker, student and office users, a family buyer, a minimalist user, an AI researcher, a creative professional, a software developer, a data analyst, a component reseller, a small-business IT buyer, a home-server operator, and a system integrator.

For each of five seeds, experiments follow three paths:

1. hold bidders at eight and vary goods from four to ten;
2. hold goods at eight and vary bidders from four to ten;
3. vary goods and bidders jointly from 4×4 to 10×10 .

The common 8×8 anchor is counted once, giving 19 environments per seed and 95 environments overall. Within a seed, successive selections are nested. Across seeds, every cell has five distinct good sets, bidder sets, and combined compositions. These five composition seeds measure sensitivity to sampling from the fixed master population; they do not estimate variation across independently generated populations or auction domains.

7.3 Experimental Comparisons

The full-information allocation is an oracle benchmark rather than a feasible communication treatment. The final analysis distinguishes:

1. the analytical full-support burden $\sum_i (2^{|G|} - 1)$;
2. the final inferred support after item exclusion and evidence-grounded `choose_one` filtering;
3. the static allocation induced by calibrated provisional XOR bids on that support;
4. the final sealed and clock allocations after deterministic exact refinement.

The mechanism comparison starts sealed and clock arms from the same frozen interest-map-plus-PV state. Policy selection is reported separately through five-seed 8×8 ablations; the selected policies are then frozen for the 95-case scalability evaluation.

7.4 Theory-Led Expectations

The following expectations were formulated to interpret the final analysis; they are not presented as preregistered hypotheses. First, excluding one good removes every bundle containing it, so the absolute support reduction from an interest map should grow rapidly with the number of goods even though the remaining candidate set can still be exponential. Second, a multiplicative PV calibration should improve outcomes only when valuation error contains a transferable systematic component; held-out payment error, rather than the PC auction result, must therefore select the scale.

Third, sealed allocation feedback should search a broader frontier and use more exact queries than a clock policy restricted to price-revealed demands. That broader search should usually improve allocation recovery, while the clock should offer a query-efficiency advantage. Fourth, selectively replacing provisional atoms with exact values need not improve hidden welfare monotonically: the WDP continues to compare corrected atoms with uncorrected ones. Finally, VCG payment reconstruction should lag allocation recovery because each winning payment requires an additional bidder-removal optimum.

Table 3: Theory-led expectations and their empirical tests.

Expectation	Rationale	Principal test
Interest maps increasingly reduce support	Exclusion removes all supersets containing an excluded item	Candidate counts along the goods path
Calibration reduces systematic PV error	A common multiplicative bias can transfer across domains	Held-out synthetic payment error
Sealed searches more broadly; clock queries less	Repeated allocation frontiers versus revealed-demand eligibility	Paired efficiency and exact VQs
Selective correction can be non-monotonic	Exact and provisional atoms coexist in the WDP	Round trajectories and worse-case counts
Payment recovery lags allocation recovery	VCG requires winner-removal welfare	Efficiency versus normalised payment error

7.5 Metrics

The main outcome is true welfare:

$$W^{true}(a) = \sum_i v_i(a_i).$$

The global full-information optimum is

$$W^* = \max_a \sum_i v_i(a_i),$$

and global efficiency is

$$\eta(a) = \frac{W^{true}(a)}{W^*}.$$

The mechanism-specific full-information outcome a_M^{FI} is also computed where appropriate, giving

$$\eta_M^{rel}(a) = \frac{W^{true}(a)}{W^{true}(a_M^{FI})}.$$

The implementation records reported welfare, true welfare, reported-space VCG revenue, full-information VCG revenue, and true surplus. Communication diagnostics include exact value queries, preparation LLM calls and tokens, PV chunks, candidate bundles, and supplementary atoms. Mechanism diagnostics include clock rounds, event usefulness, winning-bundle coverage, final-allocation hits, termination reasons, and failure classifications. Payment analysis separates allocation error from counterfactual payment error.

The reported costs must not be conflated. Table 4 separates bidder interaction from offline machine preparation. Exact-query counts measure formal preference requests selected by the mechanism, but the principal experiments answer them by deterministic lookup. They are therefore a protocol-burden measure, not direct evidence of human effort.

7.6 Candidate and Interest-Map Diagnostics

Because full hidden valuation tables are available, candidate-support quality can be measured directly. Candidate reduction is

$$1 - \frac{|\mathcal{C}_i|}{2^{|G|} - 1}.$$

Winning-bundle coverage checks whether full-information allocation bundles are present in candidate supports. Top- k coverage compares each candidate set with the bidder’s top hidden-value bundles. Interest-map diagnostics compare inferred and oracle item sets, substitute groups, acquisition modes, and resulting candidate supports.

Table 4: Information and cost layers in the experimental architecture.

Layer	Reported measure	Interpretation
Person interaction	One opening answer; number of exact VQs	Formal requests a deployed bidder would face
Initial formal support	Candidate bundle count	Bundle values the proxy must initially score
Machine preparation	LLM calls, tokens, chunks, latency	One-off frozen-pack construction cost
Auction computation	WDP solves and clock rounds	Deterministic mechanism runtime
Human cognitive burden	Not directly measured	Requires a human-subject study

The static initial-PV allocation and its full-information welfare measure the combined consequence of support selection and provisional scoring. Detailed diagnostics still distinguish omitted-bundle failure from miscalibration: a proxy may fail because it did not consider a relevant bundle, because an incorrect acquisition mode filtered it, or because it considered the bundle but assigned an inaccurate provisional value.

7.7 Frozen Specification and Reporting Protocol

The complete population and all 95 planned compositions pass deterministic population, structural, and economic preflight validation. Frozen elicitation packs were generated once per seed and goods count, then reused across every mechanism treatment. The final specification fixes the calibration scale at 1.4758; the sealed policy at competitive feedback with incumbent verification, winner-removal counterfactuals, scarcity fallbacks, and large-correction follow-up; and the clock at top-three demand, a \$50 price step, natural no-excess-demand clearance, and terminal revealed-witness/winner closure. Model-role robustness is deliberately deferred rather than mixed into the primary dataset.

Plots and tables report the mean and dispersion across five seeds for each cell rather than joining one point from each nested path as if it were an independent trend. Because the 19 cells within a seed are nested and correlated, uncertainty intervals are computed from five seed-level means using a Student- t critical value; the 95 cells are not treated as 95 independent replications. The principal tables include:

1. environment, calibration, and frozen-pack validation;
2. full-powerset and final inferred candidate support;
3. static, sealed, and clock welfare, efficiency, queries, and rounds;
4. policy ablations and event-type frequency, correction magnitude, and VCG-witness incidence;
5. reported versus full-information VCG payment diagnostics.

8 Results

The final analysis contains 95 matched auction cases: 19 scalability cells for each of five composition seeds. Every sealed-clock pair starts from the same frozen disclosure, inferred support, raw provisional values, and calibration. Unless stated otherwise, reported percentages are equally weighted case means. Confidence intervals use the five seed-level means as the independent observations.

8.1 Generated Environment Structure

The accepted master population contains 16 goods in eight functional categories and 16 bidder archetypes in four strata. Figure 2 summarises its competitive and economic structure. A bidder assigns positive base value to 78.1% of the master catalogue on average, with bidder-level density ranging from 43.8% to 100%. Goods attract between nine and 16 positive bidders (mean 12.5), so exclusions coexist with broad competition. The private profiles contain 27 `choose_one` substitute groups, 29 `can_use_multiple` groups, and 52 complement groups.

Across the 95 accepted experimental cells, every selected good has at least two positive bidders and every cell passes the preregistered structural and economic checks. The full-information optimum has a median of three winners (mean 3.0); the largest winner contributes 59.3% of optimum welfare on average and never more than 88.4%. All selected goods are allocated in the full-information optimum. These diagnostics establish that the suite contains competition and non-additive structure rather than mechanically concentrating all value in one bidder. They describe one generated PC-component population, not a representative sample of all combinatorial-auction domains.




Figure 2: Structure of the generated experimental environment. Panel (a) reports positive bidders per master-population good; panel (b) reports bidder interest density by stratum; panel (c) reports full-information winner counts across accepted cells, with colour indicating the largest winner’s welfare share and marker size indicating bidder count.

8.2 Provisional-Value Calibration

The out-of-domain calibration selects a uniform scale of 1.4758. Across nine held-out synthetic environments, mean normalised payment error falls from 0.289 without calibration to 0.233 with calibration, a 19.5% relative reduction; seven of nine held-out cases improve. The scale is therefore used for all initial provisional bids in the final PC-component experiments.

8.3 Interest-Map Support Reduction

Figure 3 compares the full non-empty powerset with final candidate support. Final support excludes every bundle containing an inferred-uninterested good and every bundle containing multiple members of an evidence-grounded `choose_one` group. Across all 95 cases it is 55.5% smaller than the full-support baseline. With eight bidders and ten goods, support falls from 8,184 bidder–bundle valuations to 2,619.2 candidates, a 68.0% reduction. At the 8×8 anchor, it falls from 2,040 to 757.6 on average.

This reduction changes the practical level of the communication problem but does not remove exponential worst-case growth. Its interpretation is also deliberately narrow. The opening disclosures are validated against hidden profiles before interest-map inference, so high




Figure 3: Auction-level candidate-support size with eight bidders as the number of goods varies. Points are five-seed means; the vertical axis is logarithmic. The full baseline is $8(2^{|G|} - 1)$.

reconstruction accuracy is a pipeline check on reconstruction from validated disclosure, not evidence of unrestricted preference recovery from arbitrary language.

8.4 Event-Policy Ablations

Policy selection uses five-seed 8×8 ablations and precedes the 95-case evaluation. In the sealed study, incumbent verification alone achieves 92.2% mean efficiency with 11.2 queries. Adding winner-removal counterfactuals raises efficiency to 95.6% with 19.0 queries; adding scarcity fallbacks reaches 97.8% with 22.4; and the full *adaptive competitive-frontier refinement* (ACFR) policy, including large-correction follow-up, reaches 99.1% with 26.2. Removing incumbent verification produces the largest degradation, to 89.4%.

The clock ablation holds the top-three, \$50-step price path fixed and varies only terminal events. PV-only supplementary allocation achieves 88.1% with no exact query. A single revealed-witness treatment reaches 90.2% with three queries; winner closure reaches 92.2% with 11.2; and the selected revealed-winner sandwich reaches 95.7% with 13.4. Full frontier closure is slightly less efficient at 95.2% while using 21.4 queries. The sandwich is therefore the best observed efficiency–query trade-off among the tested policies, rather than a claim that adding events is monotonically beneficial.

A leave-one-composition-seed-out stability check provides a limited guard against selection by one favourable ordering. Maximising mean efficiency on four seeds selects ACFR in all five sealed folds, with a tie against the policy without scarcity in one fold; ACFR’s held-out mean is 99.1%. The revealed-winner sandwich is selected in all five clock folds and retains 95.7% held-out mean efficiency. This supports stability across the available compositions, but it is not validation on a newly generated population.

8.5 Full-Suite Overview

Table 5 gives the main paired results. The calibrated initial bids achieve 86.4% mean efficiency. Sealed elicitation raises this to 94.8%, a paired gain of 8.4 percentage points (seed-clustered 95% CI 3.9–12.9), and improves 78 of 95 cases. Clock elicitation raises mean efficiency to 92.8%, a 6.5-point gain (95% CI 1.3–11.6), and improves 68 cases. Selective exact correction is




Figure 4: Five-seed 8×8 event-policy ablations. The selected policies are adaptive competitive-frontier refinement for sealed and the revealed-winner sandwich for clock.

non-monotonic: sealed worsens two cases and clock worsens 12 because the WDP continues to compare corrected atoms with remaining provisional atoms.

Table 5: Final allocation and communication results over 95 matched cases. Efficiency intervals use five seed-level means. Better/same/worse compares each elicited mechanism with its common initial provisional allocation.

Arm	Mean eff.	Median	95% CI	$\eta \geq 95\%$	Exact	Mean VQs
Initial PV	86.4%	88.2%	82.3–90.5%	19	7	0.0
Sealed	94.8%	97.8%	93.7–95.8%	57	23	21.8
Clock	92.8%	95.0%	88.8–96.8%	48	15	12.3

Figure 5 shows all three scalability paths. Mean sealed efficiency is 95.0% when bidders vary at eight goods, 94.5% when goods vary at eight bidders, and 96.0% along the joint path, with the common anchor included in each visual series. Clock gives 94.4%, 92.4%, and 92.5%, respectively. The curves are non-monotonic because successive cells are nested changes in composition rather than iid draws, but they show no collapse in average efficiency through the 10×10 endpoint.

The final event policies produce a clear query trade-off. Figure 6 shows 21.8 exact VQs per sealed case and 12.3 per clock case. Clock uses 9.6 fewer queries on average (seed-clustered 95% CI 7.4–11.7 fewer) and is less query-intensive in 86 of 95 paired cases. This is a consequence of the final clock redesign: price discovery itself makes no exact query, and only revealed terminal winners and VCG witnesses are eligible for verification. The sealed mechanism searches its allocation and counterfactual frontier during every elicitation round.

All 95 clock auctions clear through no excess demand. The median clock lasts 30 rounds, the mean 33.6, and the maximum 84. Seventeen cases require more than 50 rounds, confirming that the frozen 500-round value is a safeguard rather than a binding termination rule. Sealed elicitation stops after a median of five and a maximum of 29 refinement rounds, below its 40-round safeguard.

8.6 Sealed Outcomes and Event Path

The final suite records allocation incidence and reported-VCG-witness incidence for every exact query. In sealed ACFR, allocated bundles account for 742 queries and winner-removal counterfactuals for 683; 30.5% and 31.3%, respectively, appear in the final reported VCG witness set. Large-correction follow-ups account for 288 queries with a 27.1% witness hit rate. Scarcity




Figure 5: True allocation efficiency across the bidder, goods, and joint scalability paths. Lines show cell means over five seeds and points show seed-level observations. The 8×8 anchor appears in each panel.




Figure 6: Deterministic exact value queries across the final scalability suite. Offline disclosure, interest-map, and provisional-valuation calls are excluded.

fallbacks account for 361 queries and have a lower 11.6% direct witness rate. That lower rate does not establish zero causal value: an exact fallback can alter an intermediate allocation and expose a later winner or counterfactual.

Across all sizes, ACFR reaches 94.8% mean efficiency with 21.8 exact queries, improving 78 of 95 initial allocations. Its median of five refinement rounds reflects repeated WDP feedback rather than a fixed query schedule. In the seed-zero four-good, eight-bidder example, incumbent and counterfactual events raise efficiency from 45.7% to 96.6% in five rounds and 15 queries. This illustrates the intended sealed logic: verify the reported winner, expose its nearest allocation competitors, and re-optimize after correction.

8.7 Clock Outcomes and Event Path

In the clock mechanism, winner closure accounts for 754 queries with a 29.7% reported-witness hit rate. The staged revealed-witness phase accounts for 125 queries with a 40.8% hit rate, while the frozen single-pass revealed-witness stage is the most targeted: 57.5% of its 287 queries enter the final reported witness set. These diagnostics show how the clock’s demand path restricts verification without eliminating payment-relevant bundles.

The clock reaches 92.8% mean efficiency with 12.3 exact queries and improves 68 of 95 initial allocations. Every price path clears through zero excess demand before terminal verification. Its lower query count follows directly from price-path eligibility: an unobserved losing atom cannot enter the revealed-witness frontier. The trade-off is narrower corrective coverage and 12 cases in which terminal correction lowers efficiency relative to the initial supplementary allocation.




Figure 7: Reported VCG-witness incidence by event type over the final 95-case suites. Incidence is descriptive rather than a causal ablation.

Figure 8 contrasts the successful sealed path above with a clock failure. The seed-zero 6×6 clock case begins with a fully efficient supplementary allocation but terminal corrections select an allocation at 76.2% efficiency after 13 queries. Exact local corrections need not improve the global allocation while other atoms remain provisional.

8.8 Cross-Mechanism Comparison and the Payment-Counterfactual Gap

Sealed exceeds clock efficiency by 1.9 percentage points on average. Clock is higher in 21 cases, the mechanisms are equal in 22, and sealed is higher in 52. However, the seed-clustered 95% interval for clock minus sealed is -5.1 to 1.2 points, so five composition seeds do not resolve a population difference in efficiency. Figure 9 shows both the concentration near equality and several mechanism-specific failures.




Figure 8: Illustrative final-suite trajectories: the largest sealed efficiency gain (left) and largest clock terminal welfare reduction (right). These are diagnostic examples, not representative averages.




Figure 9: Paired sealed and clock efficiency under the same frozen initial proxy state. The dashed line denotes equal efficiency.

Payment reconstruction is a separate and more demanding information problem, as Table 6 makes explicit. An efficient allocation can be supported by accurate values for a small set of winning and displacing bundles. VCG additionally requires a correctly reported bidder-removal optimum for every winner.

Table 6: Different information requirements for allocation and VCG pricing.

Outcome	Required reported information	Sparse-elicitation failure
Efficient allocation	Values sufficient to rank feasible allocations near a^*	Winning or close competing bundle is omitted or misranked
Winner i 's VCG payment	Allocation information plus accurate W_{-i}^*	Bidder-removal witness remains provisional, omitted, or unrevealed

Empirically, payment reconstruction remains materially weaker than allocation recovery. Mean absolute per-winner payment error, normalised by optimum welfare, is 32.1% for sealed and 36.5% for clock. The corresponding absolute revenue errors are 12.0% and 13.4%. Paired payment-error and revenue-error differences both have seed-clustered intervals crossing zero. Figure 10 shows that high allocation efficiency does not force reported-space VCG payments to approach the oracle payments: each winner's payment also depends on welfare in a bidder-removal allocation whose support may remain provisional or unrevealed.




Figure 10: Absolute reported-space VCG payment error normalised by full-information optimum welfare. Lower is better.

9 Discussion

9.1 Support Selection and Provisional Scoring

The interest map and provisional valuation solve different problems. Interest-map inference is a support-selection decision: excluded items and evidence-grounded acquisition modes determine which bundles can enter the initial bid. Provisional valuation is a scoring decision conditional on that support. A candidate omitted by the interest map cannot be recovered merely by improving the PV prompt. Conversely, a broad support does not guarantee a good allocation if allocation-relevant bundles are badly ranked.

This decomposition is useful both scientifically and operationally. Exact values on interest-map support measure the best outcome available after support filtering. Comparing that benchmark with proxy PV identifies scoring loss. Comparing oracle and inferred candidate sets identifies exclusion and substitute-mode loss. It also prevents interest-map performance from being judged only through the final auction outcome, where several errors interact.

9.2 Mechanisms as Elicitation Environments

Sealed and clock auctions begin from the same frozen proxy state but expose different reasons to revisit a bundle. Competitive sealed feedback searches the reported allocation frontier directly throughout its iterative phase. The final clock policy instead uses the price path as a relevance filter: it first discovers demand without exact queries, then verifies newly exposed winners and bidder-removal witnesses only when those witnesses were revealed on the clock path. The event interface makes these differences explicit without embedding LLM orchestration inside the auction mechanism.

The ablations show that more querying is not automatically better. Full clock-frontier closure uses eight more queries than the selected sandwich yet has slightly lower mean efficiency. Direct event “hits” must nevertheless be interpreted cautiously. A queried bundle need not appear in the final or full-information allocation to be causally useful. Correcting it can change an intermediate WDP allocation, which exposes another winner or counterfactual bundle. Event evaluation therefore combines direct final-allocation incidence with trajectory and ablation evidence.

9.3 Non-Monotonic Refinement

Refinement is not guaranteed to improve true welfare monotonically. The WDP optimises a mixture of exact queried values and uncorrected provisional values. Correcting one bundle can therefore move the reported optimum to an allocation with lower hidden welfare, even though the individual correction is exact. This is a central feature of selective elicitation rather than necessarily an implementation defect. Useful query policies must be both value-correcting and allocation-relevant, and reported trajectories should not be assumed to converge monotonically.

9.4 Preparation Cost and Person-Side Query Cost

The implementation separates person-model disclosures, verification, interest-map inference, PV chunks, value queries, and demand queries. This matters because their costs have different interpretations. Person and proxy tokens are one-off offline preparation costs attached to frozen packs. Mechanism-triggered VQs are person-side preference queries, but in the principal experiment they are deterministic and consume no LLM tokens. Mechanism runtime therefore scales in WDP solves, clock rounds, and exact queries rather than hosted-model calls.

Interest-map filtering reduces the candidate powerset by 55.5% on average, but support generation still grows rapidly for broad-interest bidders. PV chunking improves structured-output reliability; it does not change the underlying number of candidate bundles that must be scored. Candidate count and preparation cost must therefore remain conceptually separate from auction-stage exact-query counts. The reported mechanism comparison concerns the latter because all offline inference is shared and amortisable.

9.5 Allocation and Payment Accuracy

Sparse elicitation places different information demands on allocation and VCG pricing. Efficient allocation requires enough information to rank feasible outcomes near the optimum. Accurate VCG payments additionally require winner-removal welfare. A large revenue difference at high allocation efficiency is therefore informative rather than contradictory: it indicates that the allocation frontier was learned more accurately than the payment counterfactual frontier. VCG-witness logging makes this distinction observable, but the residual 32–36% normalised payment error shows that the selected sparse policies do not fully close bidder-removal support.

9.6 Expected and Observed Behaviour

The results broadly support the theory-led expectations in Table 3, but with qualifications. Support reduction does increase materially with goods, although broad-interest profiles preserve a large residual candidate set. Held-out calibration reduces mean payment error in seven of nine calibration environments, supporting a systematic component without implying universal improvement. Sealed refinement is more query-intensive and has higher mean efficiency than clock, as expected, but the clustered interval on their efficiency difference crosses zero. The observed worse cases confirm the predicted non-monotonicity of partial correction. Most strongly, both mechanisms recover allocations much more accurately than VCG payments, confirming that payment-counterfactual closure is a distinct information problem. Thus the evidence supports the proposed information decomposition more clearly than it establishes a ranking between the two mechanisms.

10 Legal, Social, Ethical, and Professional Issues

The experiments use synthetic bidder profiles and make no decision about an identifiable person. Consequently, the frozen dataset is not evidence that a deployed system would satisfy legal or ethical obligations. Deployment would process preference descriptions, budgets, purchasing priorities, and auction behaviour that may be commercially sensitive or personally identifying. The relevant professional question is therefore not whether this synthetic study itself collected personal data, but what safeguards would be required before the architecture mediated real bids.

10.1 Accessibility, Autonomy, and Distributional Effects

Natural language can lower the formal barrier to expressing package preferences, which is the project’s principal social motivation. It may help participants who cannot construct XOR bids or reason over an exponential bundle space. That benefit is currently an architectural possibility, not a measured outcome: there is no human-subject evidence about comprehension, effort, disability access, or whether repeated exact questions are easier than direct bidding.

Language interfaces can also create unequal performance. Dialect, verbosity, domain vocabulary, disability-related communication patterns, and familiarity with auctions may influence what the proxy infers. A false exclusion is particularly consequential because it removes every bundle containing the item. Human–AI guidance recommends communicating system capability, making uncertainty visible, and supporting correction when the system is wrong [Amershi et al., 2019]. A deployed proxy should therefore show the person its interest map, acquisition modes, candidate omissions, provisional winning bid, and material later corrections. The person should be able to amend them before commitment and choose direct participation instead. Such controls would preserve bidder autonomy rather than treating conversational convenience as consent to an opaque economic action.

Errors have distributional effects beyond the proxy’s owner. An overstated bid can displace another bidder or alter their payment; an understated bid can leave goods inefficiently allocated. Model quality, compute budget, or access to a stronger private proxy could therefore become a source of market advantage. Equal model access, standardised disclosure of proxy use, and audit of error rates across participant groups would be necessary deployment questions. The current generated archetypes are economic roles, not a fairness evaluation over protected characteristics.

10.2 Privacy, Confidentiality, and Security

The ICO’s AI guidance organises data protection around lawfulness, fairness and transparency, purpose limitation, data minimisation, accuracy, storage limitation, security, and accountability [Information Commissioner’s Office]. Applying those principles would require a documented

lawful basis, a clear explanation of what information is sent to each model provider, collection of no more detail than elicitation requires, retention and deletion rules, access controls, and appropriate handling of international processing. A budget or intended purchase could reveal financial circumstances or confidential business strategy even when it is not a conventional identity field.

The implementation provides useful technical controls: model roles receive only the information they require; frozen public packs exclude hidden valuation tables; cache entries and calls are logged; schemas reject malformed outputs; and deterministic replay avoids repeatedly transmitting disclosures. These controls support minimisation and auditability but do not themselves establish regulatory compliance. Real deployment would additionally require provider contracts and retention settings, encryption and credential management, incident response, user access and deletion procedures, and a data-protection impact assessment where the processing risk warrants one.

The language interface also enlarges the attack surface. A catalogue or user answer could contain prompt injection; a compromised cache could expose preferences; and an auction participant might intentionally manipulate a proxy into misclassifying alternatives. Fail-closed parsing protects the formal schema, not the semantic truth of an adversarial answer. Provider-side content isolation, authenticated catalogue data, least-privilege access, signed frozen artifacts, semantic anomaly tests, and manual review for high-value auctions would therefore be appropriate safeguards.

10.3 Economic Trustworthiness and Professional Responsibility

The BCS Code requires attention to privacy, security, wellbeing, the environment, third-party rights, non-discrimination, equal access, competence, and honest reporting of system performance [BCS, [The Chartered Institute for IT](#)]. The ACM Code similarly emphasises avoiding harm, honesty, fairness, privacy, security, and thorough evaluation of system risks [[Association for Computing Machinery, 2018](#)]. These duties are relevant because the proxy performs an economically consequential translation rather than merely generating advisory text.

The paper accordingly reports negative as well as positive results: selective correction sometimes reduces welfare, payment errors remain substantial, and the five seeds do not represent independently generated populations. A professional deployment should not advertise the observed 94.8% or 92.8% mean efficiency as a guarantee for another domain. It should expose when a bundle remains provisional, distinguish allocation confidence from payment confidence, retain an appeal or correction path, and assign responsibility for loss caused by model, software, or operator error. Strategic behaviour is also unresolved. A proxy that is optimised to win rather than report faithfully could amplify manipulation, while deterministic exact lookups in this study assume away that conflict.

Intellectual-property and sustainability issues are smaller but not absent. Deployment must respect model and dependency licences, catalogue rights, and ownership of generated profiles and transaction records. Hosted inference also has financial and environmental costs; prior work shows why computation and energy should be reported rather than treated as free [[Strubell et al., 2019](#)]. This implementation records calls, tokens, latency, and cache hits, and frozen replay removes auction-stage model calls. It does not measure electricity or carbon, so no environmental reduction is claimed. A production assessment should compare the one-off inference cost with the human and computational costs of alternative auction interfaces.

Taken together, the responsible interpretation is conditional. The framework provides inspectable information boundaries, deterministic replay, correction hooks, and audit records that could support a trustworthy service. It does not yet demonstrate inclusive usability, legal compliance, strategic safety, or acceptable payment risk. Those are deployment requirements and evaluation tasks, not consequences that follow automatically from high simulated welfare.

11 Limitations

Four limitations most directly bound the paper’s claims:

1. **One domain and one generated population.** The five composition seeds are nested samples from one validated 16-by-16 master population, not five independently generated populations. Statistical intervals therefore describe composition sensitivity within this population.
2. **No strategic behaviour.** Provisional estimates and exact refinements are treated as reports to be corrected, without bid shading, manipulation, or equilibrium analysis. The mechanisms are experimental elicitation environments rather than deployment-ready auction designs.
3. **No human-subject evidence.** The architecture reduces formal reporting and presents a natural-language interface, but it does not establish that real bidders experience lower cognitive burden or can answer repeated exact value queries consistently.
4. **Policy selection within the same population.** The 8×8 ablations and 95-case evaluation share the master population. Leave-one-seed-out stability helps only across compositions; it does not establish transfer to a new domain or independently generated population.

Additional modelling limitations remain. Population constraints and the deterministic valuation formula encode judgement about complements, substitutes, budgets, caps, and saturation. The fixed model assignment, validated disclosures, and closed-world interest map may overstate inference reliability relative to incomplete or strategic human language. Candidate support is frozen rather than revised online. The CCA-inspired clock omits activity rules, a full supplementary-bidding phase, optimised increments, and core-selecting payments. The current environment also assumes unit-supply indivisible goods. Model-role robustness, other mechanisms such as CECA, and independently generated domains remain future validation rather than evidence claimed here.

12 Extensions

Natural extensions include:

- dynamic interest-map updates after mechanism feedback;
- uncertainty-aware event priority based on pivotality and confidence;
- a bounded post-allocation payment-counterfactual audit;
- demand-query and noisy human-response robustness treatments;
- interval-valued or uncertainty-calibrated provisional valuations;
- CCA-lite activity rules and core-selecting payments;
- strategic proxy policies;
- independently generated populations, domains, and multi-model role robustness;
- human-subject validation of the person-simulator assumptions;
- CECA-style satisfaction and counterexample event adapters.

13 Conclusion

This paper develops a reproducible framework for scalable preference elicitation with LLM proxy bidders. A constrained generation-and-validation pipeline supplies structured private benchmark environments. A language-to-bid pipeline separates candidate selection through an interest map from candidate scoring through provisional valuations. A modular event interface then embeds the same frozen proxy state in iterative sealed and ascending clock auctions.

The separation between support, scoring, and refinement is central. It makes candidate omission distinguishable from provisional misvaluation and makes mechanism feedback comparable without repeated LLM sampling. The closed-world interest map provides scalability, while evidence-grounded acquisition modes protect multi-use bidders from overly aggressive substitute filtering. Frozen replay and deterministic value queries isolate the economic query path from hosted-model variability.

The framework does not remove the exponential communication problem for general combinatorial values. Instead, it provides a controlled empirical method for studying where information is discarded, how initial bids are formed, and which allocation- or price-derived events recover decision-relevant values. This makes the interaction between language inference and auction design observable, reproducible, and amenable to ablation.

In the final five-composition-seed evaluation, interest maps reduce candidate support by 55.5% relative to full powersets. Calibrated provisional bids achieve 86.4% mean efficiency; sealed and clock elicitation raise this to 94.8% and 92.8% with 21.8 and 12.3 exact value queries, respectively. The selected clock policy preserves genuine ascending-price discovery and clears naturally in all 95 cases, while its revealed-demand restriction substantially reduces query volume. Neither policy reconstructs VCG payments as successfully as it reconstructs allocations, identifying sparse payment-counterfactual closure as the most important remaining mechanism-design problem.

A Reproducibility and Additional Specification

A.1 Frozen Model Roles and Information Boundaries

The accepted master population was generated by OpenAI `gpt-5.6-sol`. Frozen opening answers and their independent verifier calls also use `gpt-5.6-sol`. Interest maps and provisional valuations use Google `gemini-3.6-flash`. The bidder proxy receives only the public catalogue and accepted opening exchange; it never receives latent profiles or complete valuation tables. All auction-stage exact answers are deterministic lookups and therefore use no model provider.

The five composition seeds $0, \dots, 4$ instantiate distinct nested goods and bidder orderings from the same master population. Each goods-specific master elicitation pack is projected to smaller bidder sets, ensuring that changing the number of rivals does not regenerate a bidder’s opening exchange or provisional values.

A.2 Frozen Preparation and Mechanism Parameters

Structured-output parsing is fail closed in the principal preparation. A failed interest map or provisional-value chunk is retried with diagnostic context and then aborts rather than silently expanding to all goods or replacing the treatment with mass exact queries. Successful model responses are committed independently to a read-write SQLite cache keyed by the rendered prompt and response-relevant settings. Frozen packs record raw and parsed responses, model role, token counts, latency, cache status, candidate support, and configuration hashes.

Table 7: Principal frozen parameters. Zero query caps denote disabled caps.

Component	Setting	Value
Opening interaction	Question policy	Canonical, one accepted answer
Interest map	World assumption / failure	Closed world / fail closed
Candidate support	Maximum bundles	Uncapped
Provisional valuation	Chunk / failure	64 bundles / fail closed
Parsing	Additional retries	2
Calibration	Family / scale / size factor	Uniform / 1.4758 / 1.0
Sealed	Events	Incumbent, winner removal, scarcity, correction follow-up
Sealed	Stop / safeguard	No new refinements / 40 rounds
Sealed	Correction / scarcity threshold	25% / 60% of incumbent report
Clock	Demand / price increment	Top three / \$50
Clock	Stop / safeguard	No excess demand / 500 rounds
Clock	Terminal policy	Single-pass revealed witness and staged winner/witness closure
Queries	Per-bidder / global cap	0 / 0
Queries	Exact response mode	Deterministic hidden-value lookup

A.3 Frozen Artifacts and Offline Reconstruction

The experiment is anchored by `outputs/final/final_experiment_v3.json`, which records scenario and pack hashes, calibration provenance, event policy, mechanism parameters, and all 95 pack paths. Raw sealed and clock outputs are stored under `outputs/final/scalability_v2` and `outputs/final/clock_sandwich_scalability_v3`. The final analysis is reconstructed without LLM calls by:

```
./venv/bin/python scripts/build_final_analytic_package.py \
--spec outputs/final/final_experiment_v3.json \
--sealed-dir outputs/final/scalability_v2 \
--clock-dir outputs/final/clock_sandwich_scalability_v3 \
--interest-map-dir outputs/interest_map_analysis \
--sealed-ablation outputs/final/event_policy_8x8 \
--clock-ablation \
  outputs/final/clock_focused_closure_ablation_pc_8x8_calibrated \
--output-dir outputs/final/analytic_package_v3
```

Generated-environment tables and Figure 2 are reconstructed from the accepted population and its validation report:

```
./venv/bin/python scripts/analyze_generated_environment.py \
--scenario-spec \
  scenarios/pc_build_v3/pc_build_population_16x16.json \
--validation-report \
  scenarios/pc_build_v3/pc_build_population_16x16.validation.json \
--output-dir outputs/final/analytic_package_v3
```

The analytic package contains case-level inputs, seed-clustered summaries, event and VCG-witness diagnostics, calibration evidence, selected trajectories, environment tables, and both

PDF and PNG figures. Rebuilding these outputs does not mutate the frozen scenario, elicitation packs, or auction runs.

A.4 Software Test Suite

The accompanying implementation has an offline regression suite of 880 tests, run from the repository root by:

```
./venv/bin/python -m pytest -q
```

The suite is organised around five verification boundaries. First, auction primitives test winner determination, deterministic tie breaking, supplementary bids, sealed VCG payments, winner-removal witnesses, clock demand, price updates, and natural termination. Second, proxy tests exercise event dispatch, sealed and clock trajectories, exact-value deduplication, candidate-bundle construction, interest-map exclusions and acquisition modes, provisional-value chunking, calibration, and query accounting. These tests include adverse cases such as malformed structured output, omitted goods, false exclusivity, incomplete chunks, repeated events, and exhausted retries.

Third, model-interface tests cover role-specific clients, prompt construction, response schemas and parsers, token and call logging, retry classification, and SQLite cache semantics. They use deterministic mock clients and temporary stores; running the test suite neither requires API credentials nor makes live model calls. Fourth, environment and preparation tests cover scenario-schema validation, generated-population constraints, deterministic nested sampling, frozen-pack projection, hidden-information boundaries, and environment fingerprints. In particular, a frozen elicitation pack is rejected if replayed against a different environment, and its serialised public data is checked not to contain the hidden valuation table.

Finally, experiment-level tests verify the scalability grid, frozen experiment specification, input hashes, resolved mechanism policies, output schemas, interest-map and scalability summaries, calibration provenance, and analytic pipeline commands. The full-information proxy is retained only as a test oracle for mechanism comparisons; it is not an experimental proxy treatment.

These are tests of implementation correctness and reproducibility, not statistical evidence for the behavioural quality of model-generated answers or for the economic performance of the mechanisms. Those claims instead rely on the frozen empirical evaluation reported in Section 8. The separation is deliberate: mock-based tests establish that a specified response is processed correctly, frozen hashes establish that the reported inputs have not changed, and the multi-seed experiments measure the consequences of those inputs.

References

- Saleema Amershi, Dan Weld, Mihaela Vorvoreanu, Adam Fourney, Besmira Nushi, Penny Collisson, Jina Suh, Shamsi Iqbal, Paul N. Bennett, Kori Inkpen, Jaime Teevan, Ruth Kikin-Gil, and Eric Horvitz. Guidelines for human–ai interaction. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*, pages 1–13, 2019. doi: 10.1145/3290605.3300233.
- Association for Computing Machinery. ACM code of ethics and professional conduct. <https://www.acm.org/code-of-ethics>, 2018.
- Lawrence M. Ausubel and Paul Milgrom. Ascending proxy auctions. In Peter Cramton, Yoav Shoham, and Richard Steinberg, editors, *Combinatorial Auctions*, chapter 3. MIT Press, 2006.
- Lawrence M. Ausubel, Peter Cramton, and Paul Milgrom. The clock-proxy auction: A practical combinatorial auction design. In Peter Cramton, Yoav Shoham, and Richard Steinberg, editors, *Combinatorial Auctions*, chapter 5, pages 115–138. MIT Press, 2006.

- BCS, The Chartered Institute for IT. BCS code of conduct. <https://www.bcs.org/membership-and-registrations/become-a-member/bcs-code-of-conduct/>. Accessed 3 August 2026.
- Gianluca Brero, Benjamin Lubin, and Sven Seuken. Machine learning-powered iterative combinatorial auctions, 2019.
- Robert Day and Paul Milgrom. Core-selecting package auctions. *International Journal of Game Theory*, 36(3–4):393–407, 2008. doi: 10.1007/s00182-007-0100-7.
- Paul W. Goldberg, Edwin Lock, and Francisco Marmolejo-Cossío. Learning strong substitutes demand via queries. *ACM Transactions on Economics and Computation*, 2022. arXiv:2005.01496.
- Theodore Groves. Incentives in teams. *Econometrica*, 41(4):617–631, 1973. doi: 10.2307/1914085.
- David Huang, Francisco Marmolejo-Cossío, Edwin Lock, and David C. Parkes. Accelerated preference elicitation with llm-based proxies, 2025.
- Information Commissioner’s Office. Guidance on AI and data protection. <https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/artificial-intelligence/guidance-on-ai-and-data-protection/>. Accessed 3 August 2026.
- Sébastien M. Lahaie and David C. Parkes. Applying learning algorithms to preference elicitation. In *Proceedings of the 5th ACM Conference on Electronic Commerce*, pages 180–188, 2004.
- Benjamin Lubin, Sven Seuken, Manuel Beyeler, and Gianluca Brero. imlca: Machine learning-powered iterative combinatorial auctions with interval bidding, 2020.
- Noam Nisan and Ilya Segal. The communication requirements of efficient allocations and supporting prices. *Journal of Economic Theory*, 129(1):192–224, 2006.
- David C. Parkes. Iterative combinatorial auctions. In Peter Cramton, Yoav Shoham, and Richard Steinberg, editors, *Combinatorial Auctions*. MIT Press, 2006.
- David C. Parkes and Lyle H. Ungar. ibundle: An efficient ascending price bundle auction. In *Proceedings of the 1st ACM Conference on Electronic Commerce*, 2000.
- Ermis Soumalias, Jakob Weissteiner, Jakob Heiss, and Sven Seuken. Machine learning-powered combinatorial clock auction, 2023.
- Ermis Soumalias, Jakob Heiss, and Jakob Weissteiner. Prices, bids, values: One ml-powered combinatorial auction to rule them all, 2024.
- Ermis Soumalias, Yanchen Jiang, Kehang Zhu, Michael Curry, Sven Seuken, and David C. Parkes. Llm-powered preference elicitation in combinatorial assignment, 2025.
- Emma Strubell, Ananya Ganesh, and Andrew McCallum. Energy and policy considerations for deep learning in NLP. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 3645–3650, 2019. doi: 10.18653/v1/P19-1355.
- William Vickrey. Counterspeculation, auctions, and competitive sealed tenders. *The Journal of Finance*, 16(1):8–37, 1961. doi: 10.1111/j.1540-6261.1961.tb02789.x.